



# CONCURRENT TCP SERVER

---

Tran Nguyen Ngoc, SoICT, HUST

# Content

- I/O Models
- Concurrent TCP server: one child per client
- Concurrent TCP server: one thread per client



# I/O MODELS

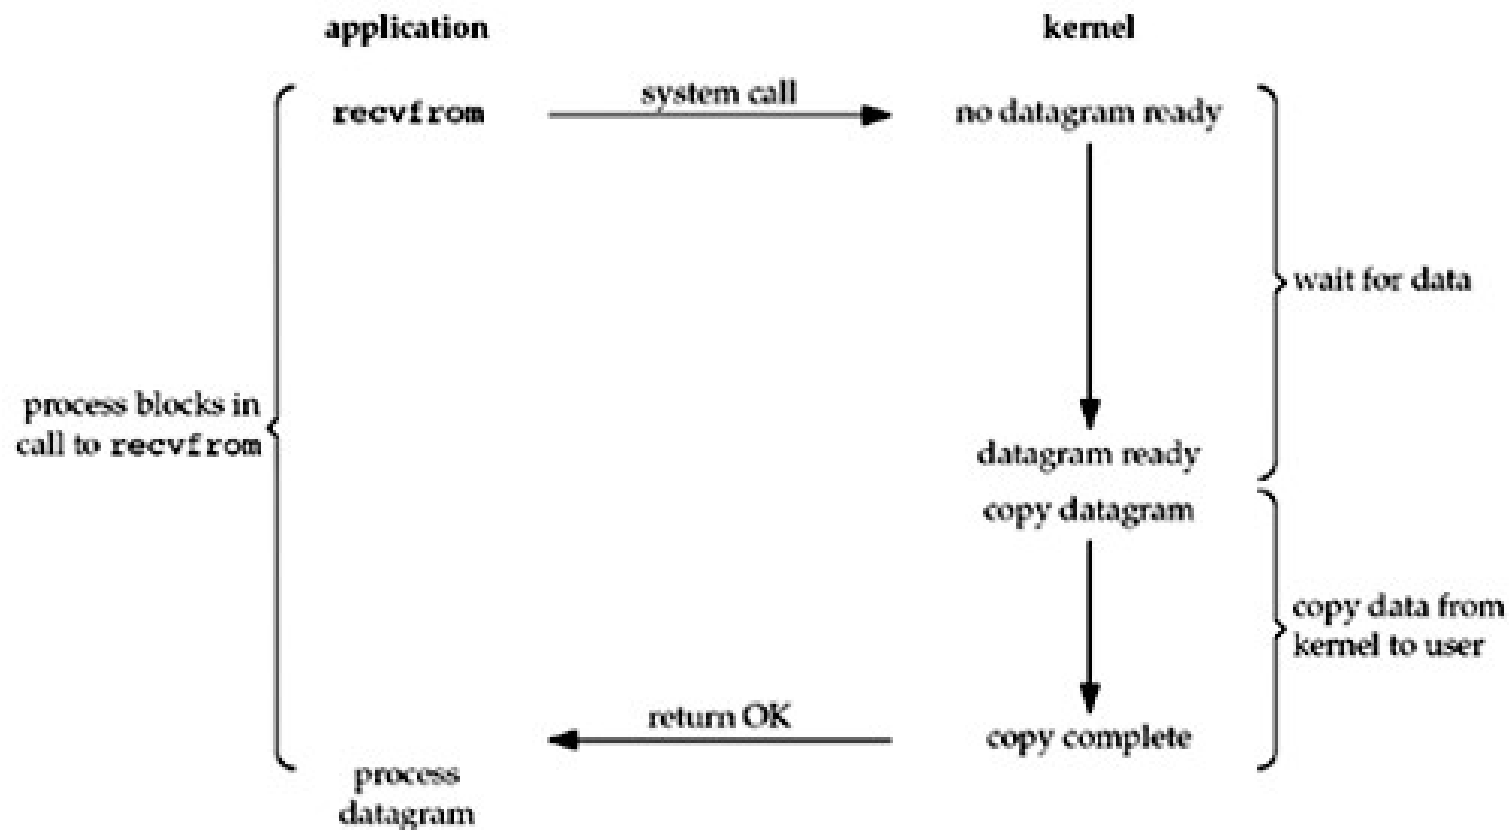
---

# I/O Models

- blocking I/O
- nonblocking I/O
- I/O multiplexing (select and poll)
- signal driven I/O (SIGIO)
- asynchronous I/O (the POSIX aio\_functions)

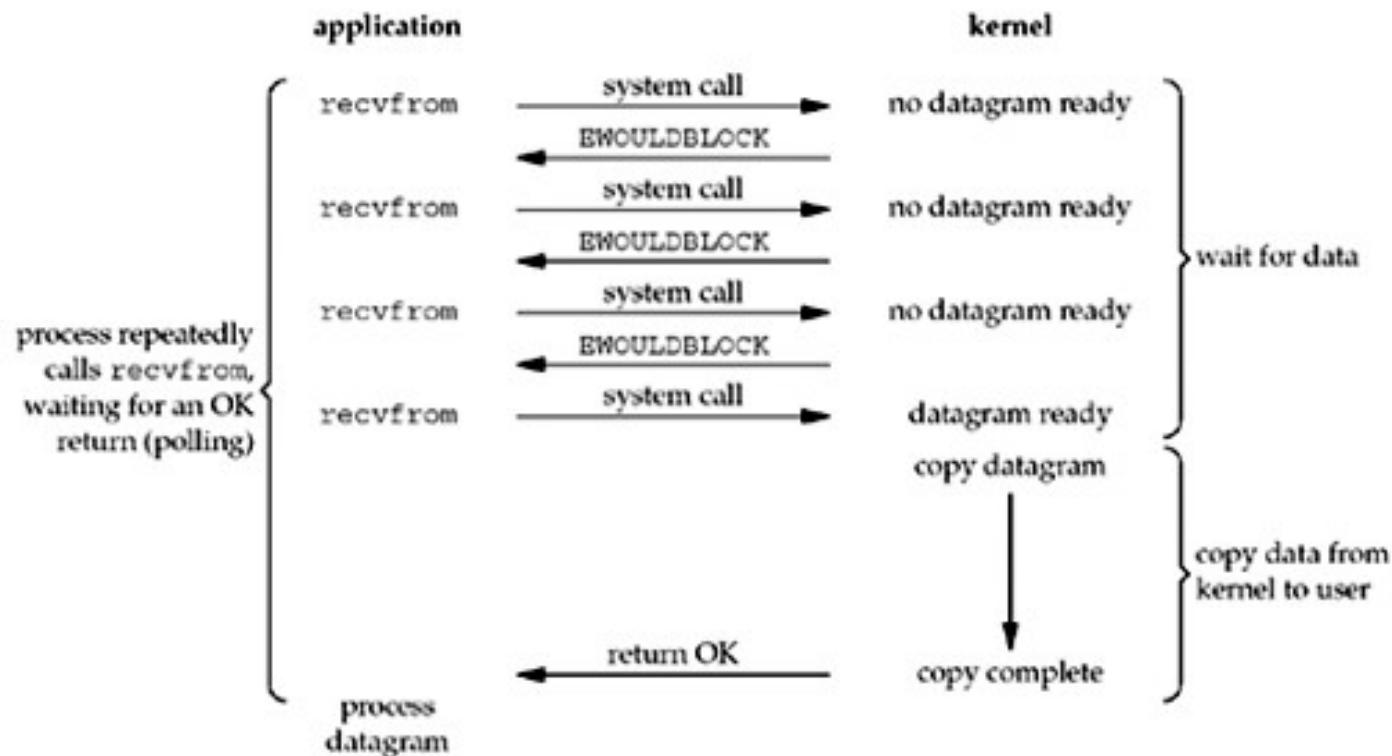
# Blocking I/O Model

- Blocking I/O model: I/O function block process/thread until returning.
- `accept()` , `connect()` , `send()` , `recv()` , ...



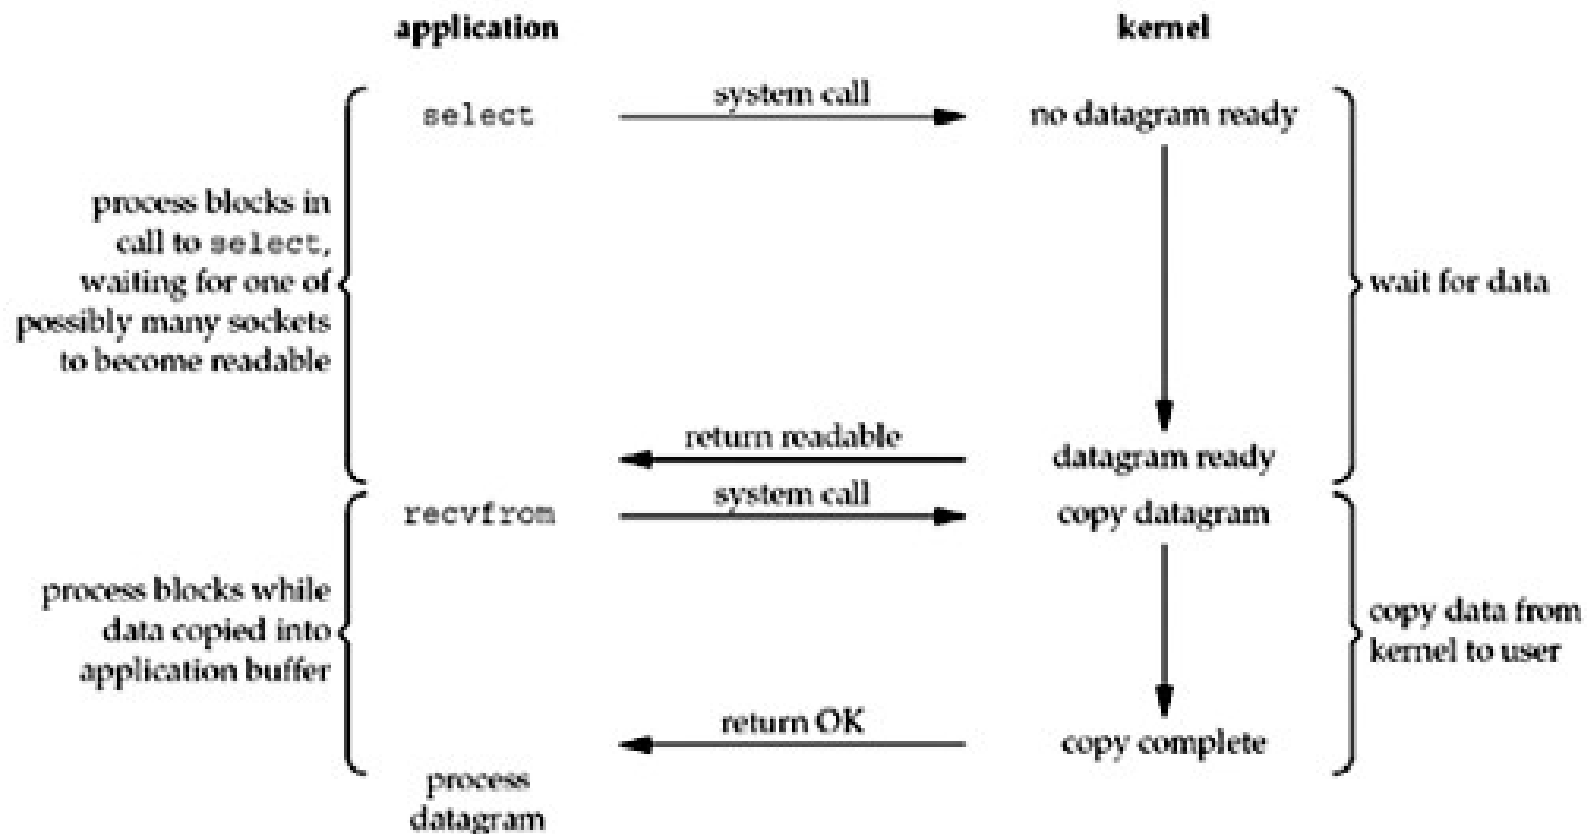
# Non-blocking I/O Model

- Non-blocking I/O model: I/O function returns immediately
- If there is no data to return, so the kernel immediately returns an error of EWOULDBLOCK instead



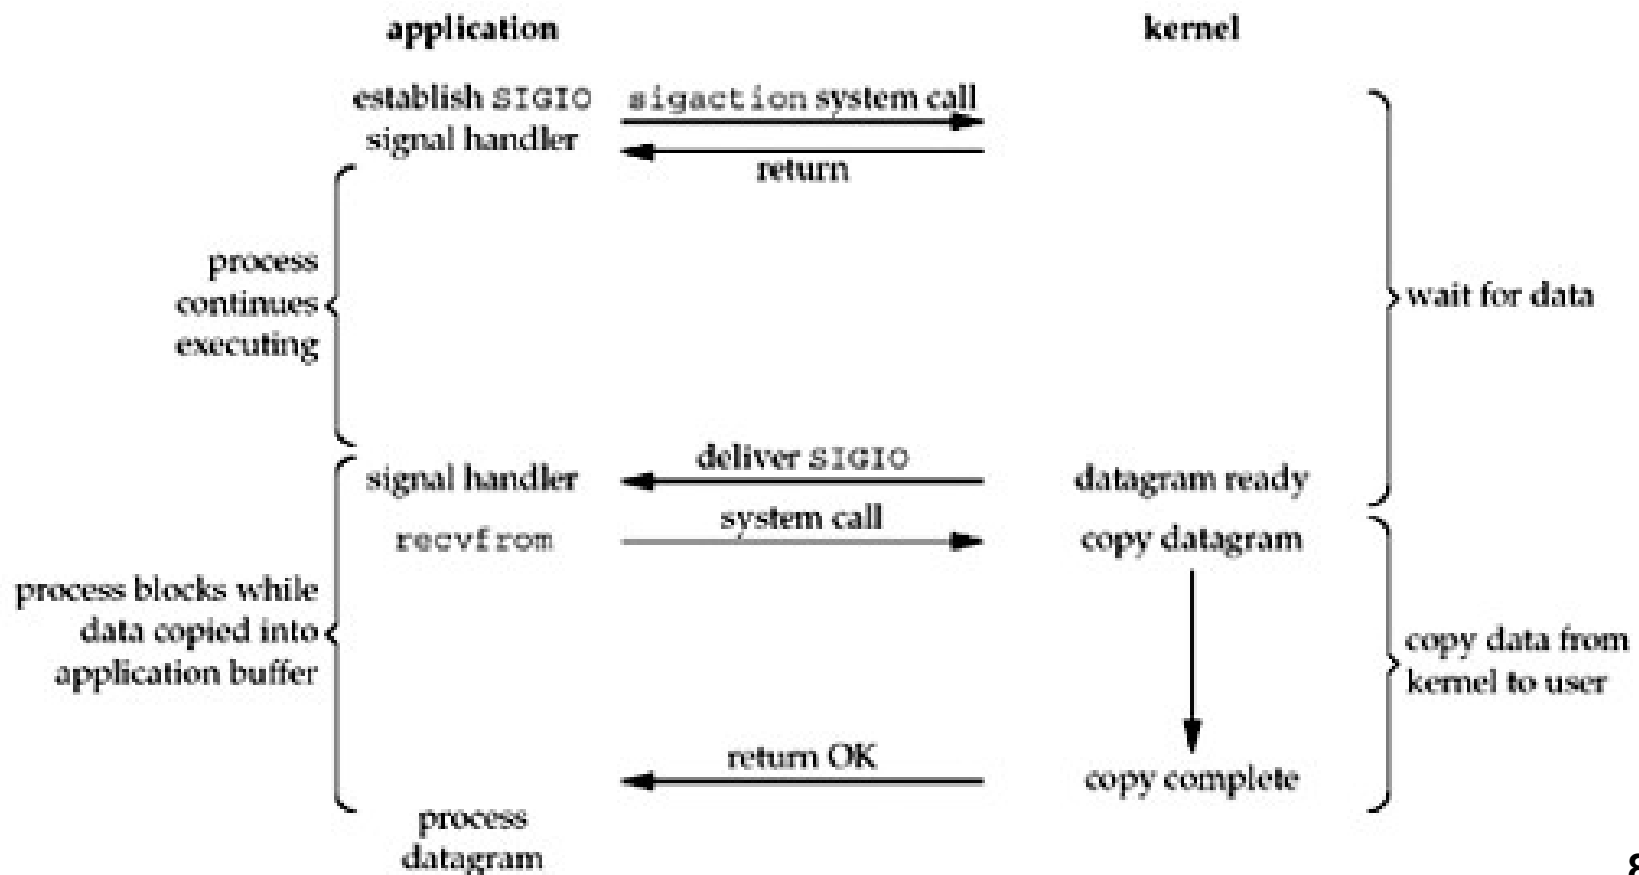
# I/O Multiplexing Model

- With I/O multiplexing, we call `select` or `poll` and block in one of these two system calls, instead of blocking in the actual I/O system call



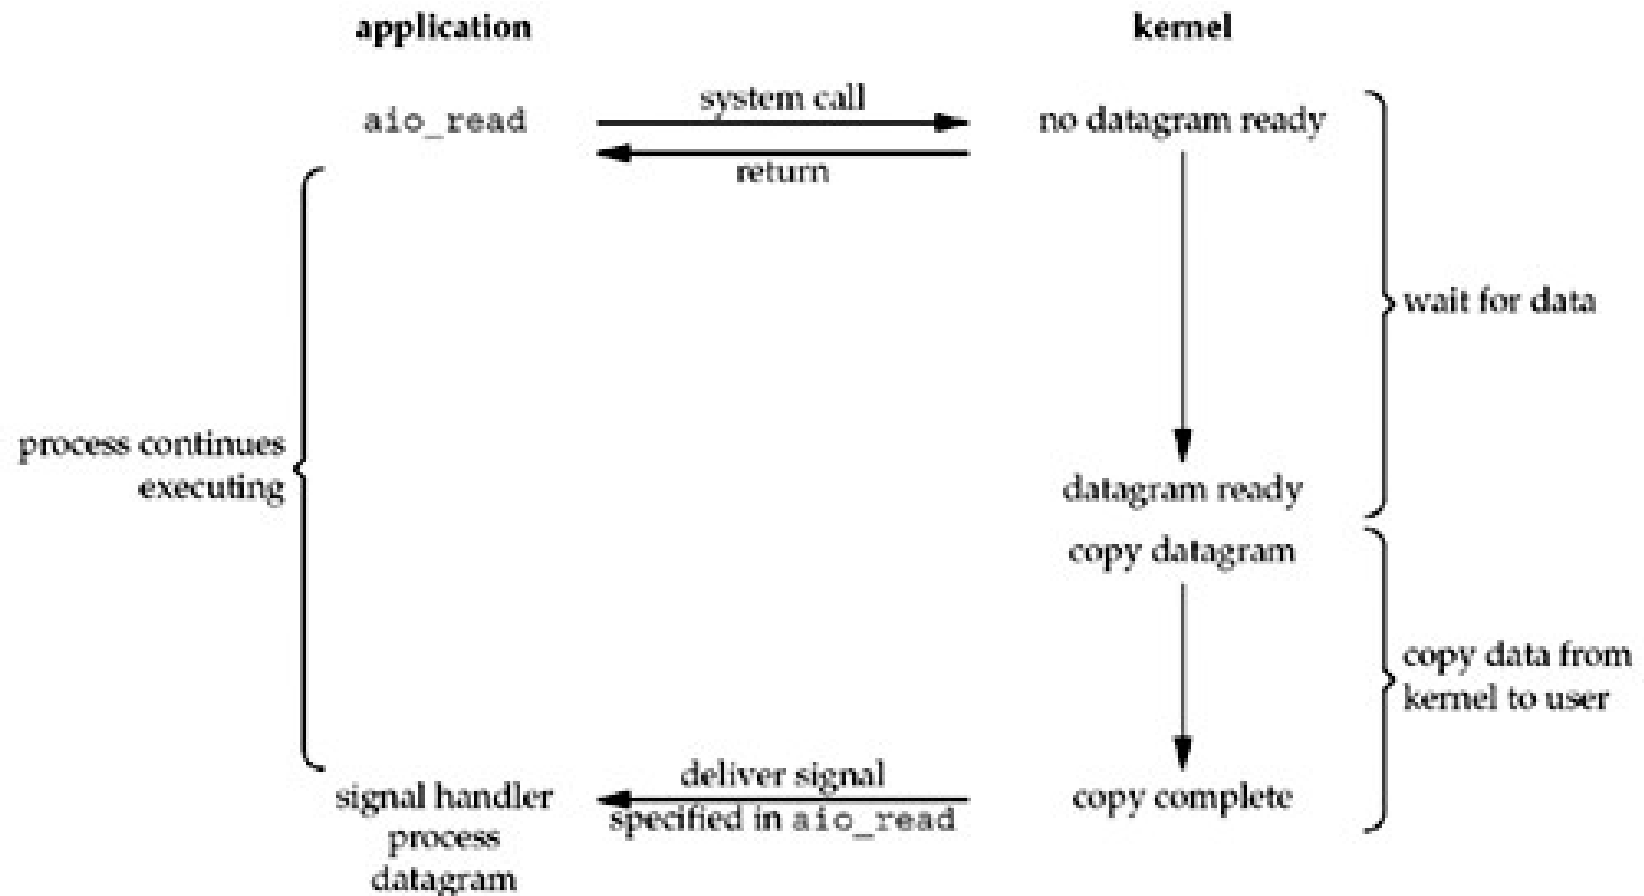
# Signal-Driven I/O Model

- Use signals, telling the kernel to notify app with the SIGIO signal when the descriptor is ready





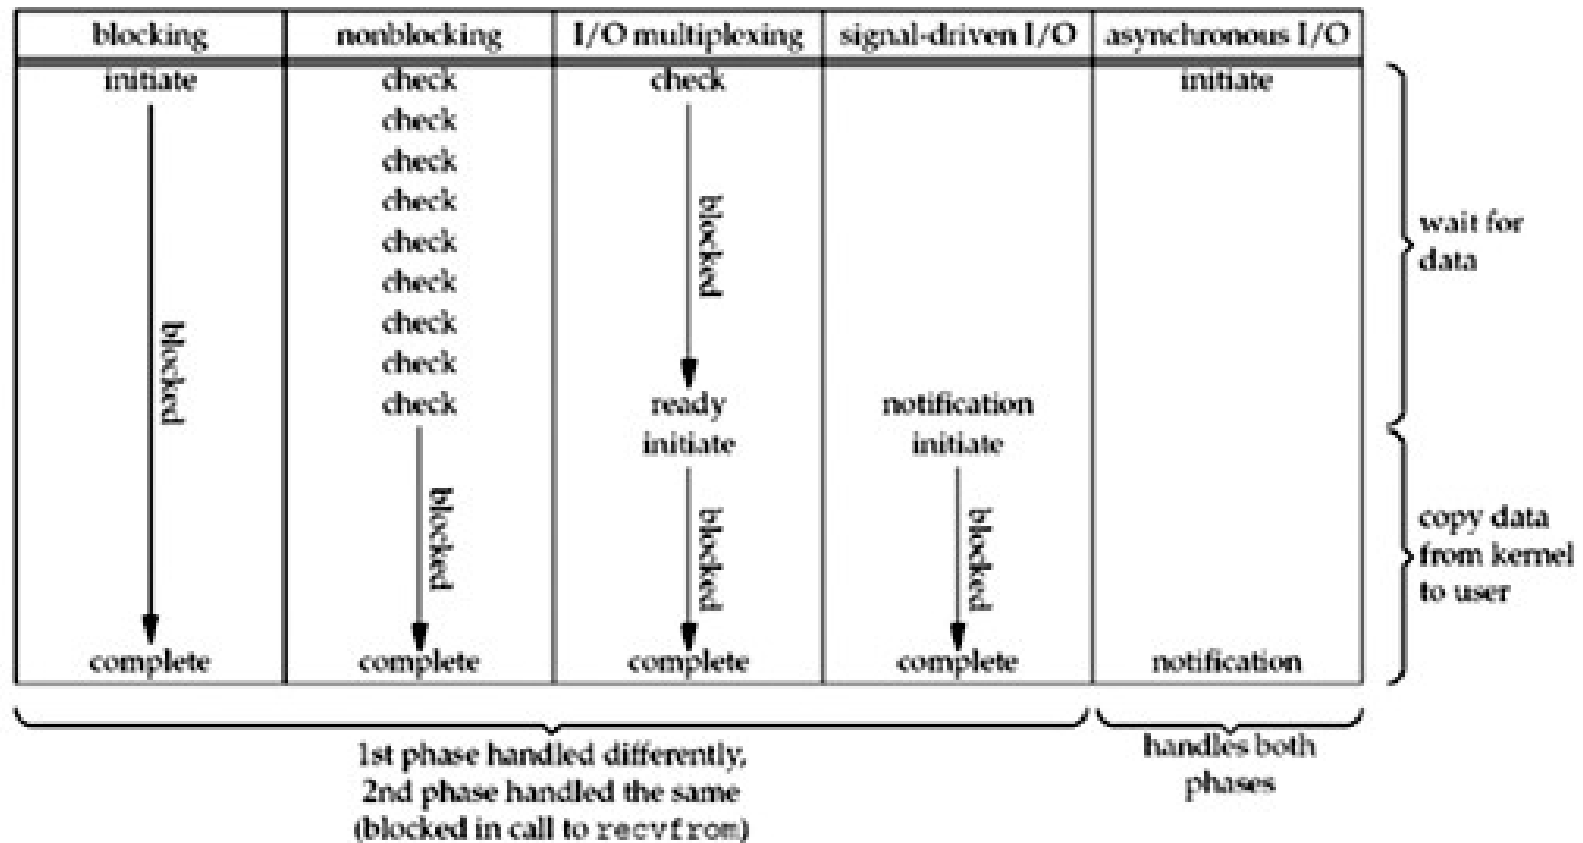
# Asynchronous I/O Model



## Asynchronous I/O Model (2)

- App calls `aio_read` (the POSIX asynchronous I/O functions begin with `aio_`)
- Pass the kernel
  - the descriptor
  - buffer pointer
  - buffer size (the same three arguments for `read`)
  - buffer offset (similar to `lseek`)
  - how to notify us when the entire operation is complete
- This system call returns immediately and our process is not blocked while waiting for the I/O to complete.

# Comparison of the I/O Models

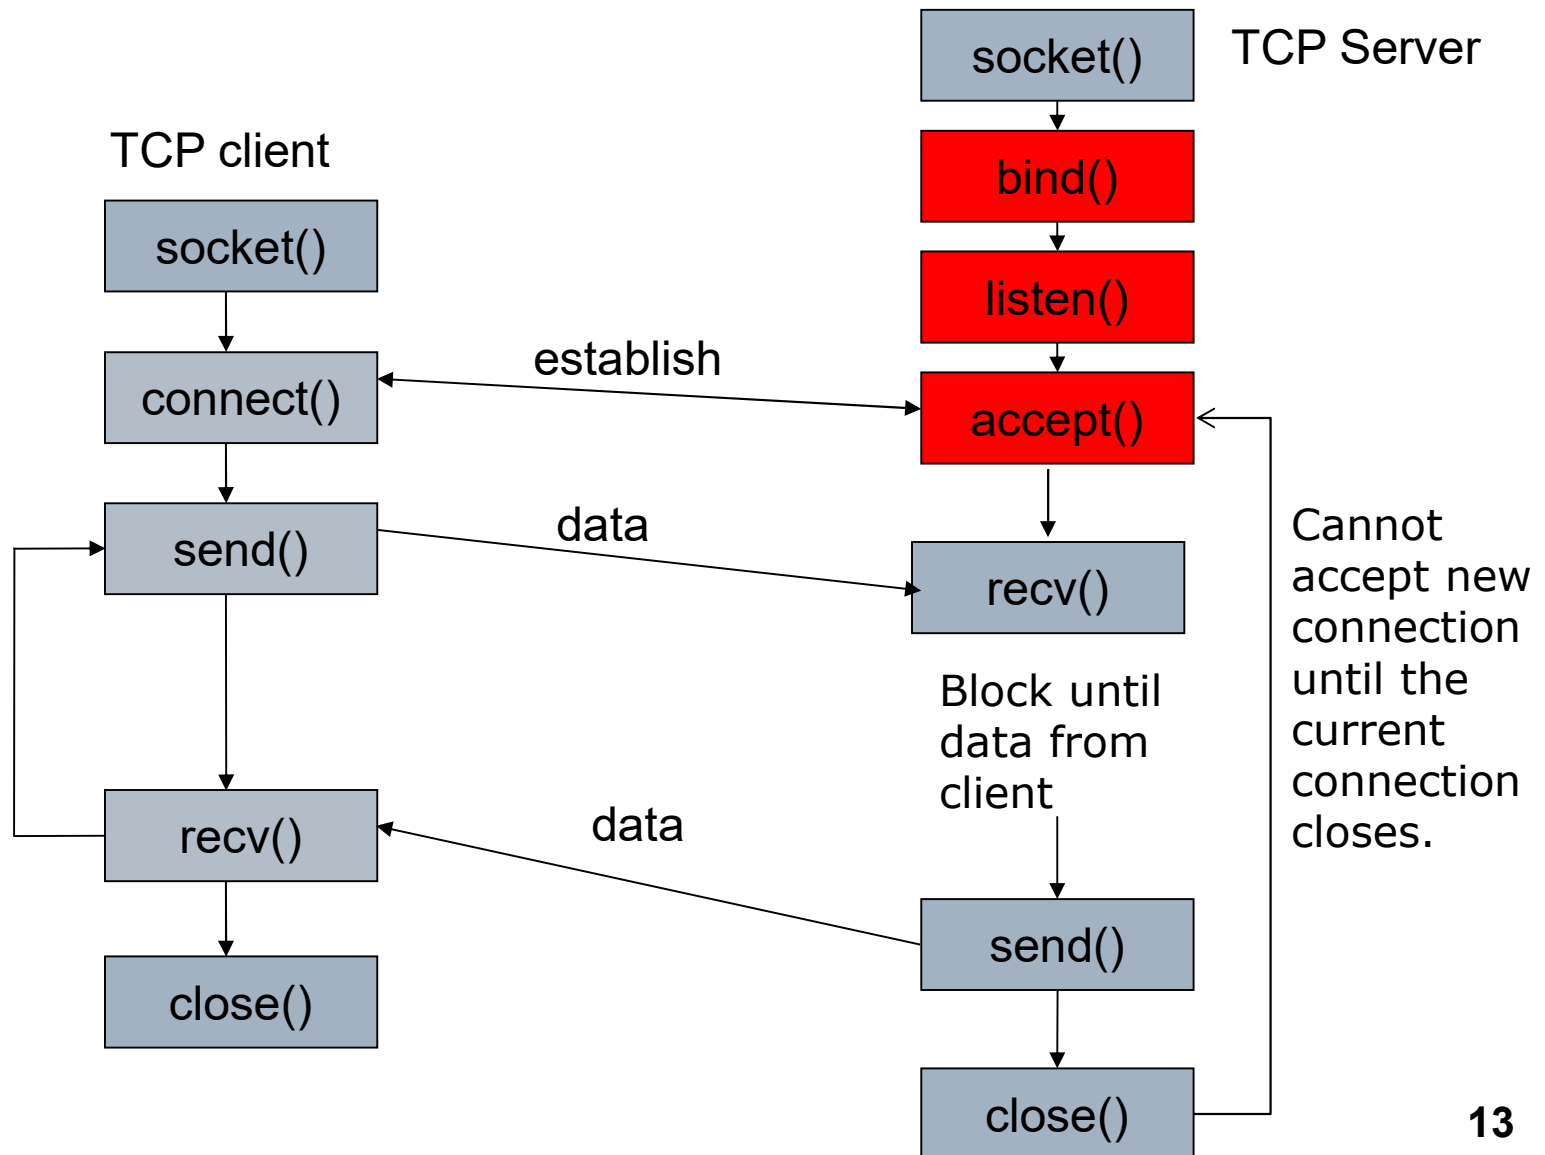


# Review TCP Echo Server

```
while(1) {
    //accept request
    connfd = accept(listenfd, (sockaddr *) & clientAddr,
                    &clientAddrLen);

    //receive message from client
    rcvBytes = recv(connfd, buff, BUFF_SIZE, 0);
    if(rcvBytes < 0) {
        perror("Error :");
    }
    else{
        buff[rcvBytes] = '\\0';
        printf("Receive from client: %s\\n",buff);
        //Echo to client
        sendBytes = send(connfd, buff, strlen(buff), 0);
        if(sendBytes < 0)
            perror("Error: ",);
    }
    closesocket(connfd);
} //end while
```

# Blocking I/O Model and TCP server



## *Iterating server*

- Simple server
  - But when a client request can take longer to service, we can't handle other clients
- Use a *concurrent server*
- One child per client: `fork()` a child process to handle each client
  - One thread per client: create a thread to handle each client by using `pthread_create()`



# MULTI-PROCESS SERVER

---


# fork()

```
#include <unistd.h>
pid_t fork(void);
```


- Create a new process by copying itself.
- Returns twice:
  - Once in the calling process (called the parent) with a return value that is the process ID of the newly created process (the child).
  - Once in the child, with a return value of 0
- All descriptors open in the parent before the call to fork are shared with the child after fork returns



# One child per client



# One child per client



# Use fork()

```
pid_t pid;
int listenfd, connfd;
//Step 1: Construct socket
//Step 2: Bind address to socket
//Step 3: Listen request from client

//Step 4: Communicate with client
while (1) {
    connfd = accept (listenfd, ... );
    if( (pid = fork()) == 0) { // process in child
        close(listenfd); // child closes listening socket
        doit(connfd);    // process the request
        close(connfd);   // done with this client
        exit(0);         // child terminates
    }
    close(connfd); // parent closes connected socket
}
```

# Handling SIGCHLD Signals

- When a child process ends, it sends the *SIGCHLD* signal to the parent
  - Information about the child process is still maintained in “process table” in order to allow its parent to read the child exit status afterward.
- If we ignore the *SIGCHLD*, the child process will enter the zombie state
- We need to wait and handle *SIGCHLD* signal

# Signaling

- A signal is a notification to a process that an event has occurred.
- Signals are sometimes called software interrupts.
- Signals usually occur asynchronously. By this we mean that a process doesn't know ahead of time exactly when a signal will occur.
- Signals can be sent
  - By one process to another process (or to itself)
  - By the kernel to a process

## Signal (cont.)

- Typing certain key combinations at the controlling terminal of a running process causes the system to send it certain signals:
  - Ctrl-C sends an INT signal ("interrupt", SIGINT)
  - Ctrl-Z sends a TSTP signal ("terminal stop", SIGTSTP)
  - Ctrl-\ sends a QUIT signal (SIGQUIT)
- SIGHUP is sent to a process when its controlling terminal is closed (a hangup)
- SIGTERM is sent to a process to request its termination.
  - Unlike the SIGKILL signal, it can be caught and interpreted or ignored by the process.

# Handling SIGCHLD Signals

- The purpose of the zombie state is to maintain information about the child for the parent to fetch at some later time.
  - They take up space in the kernel and eventually we can run out of processes
- Whenever we *fork* children, we must *wait* for them to prevent them from becoming zombies → establish a signal handler to catch *SIGCHLD*, and within the handler, we call *wait*
- Establish the signal handler by adding the function call :  
*signal (SIGCHLD, handler);*

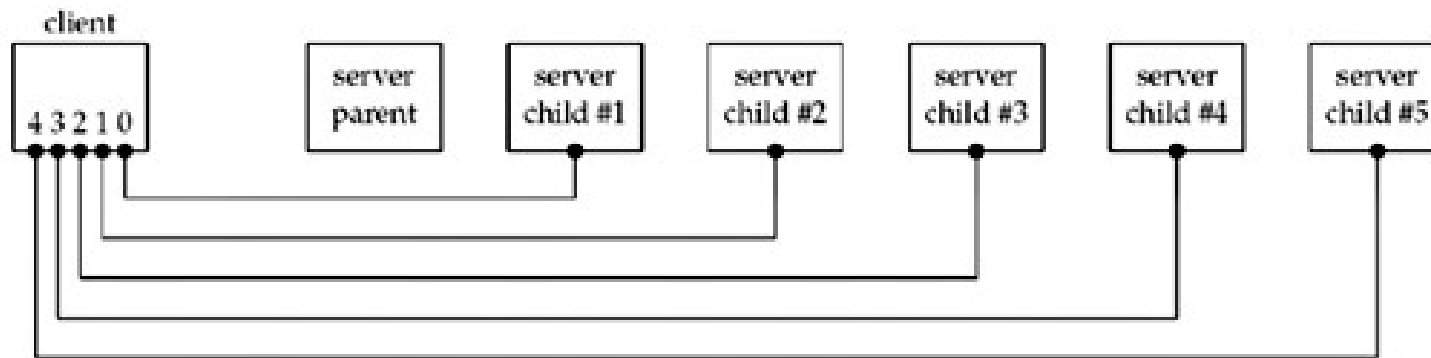
## wait() and waitpid()

```
#include <sys/wait.h>
pid_t wait (int *statloc);
pid_t waitpid (pid_t pid, int *statloc, int options);
```

- Wait for the status change of a process.
- Use to handle the terminated child
- Both return two values:
  - The return value of the function:
    - the process ID of the terminated child
    - 0 or -1 if error
  - The termination status of the child (an integer) is returned through the *statloc* pointer.

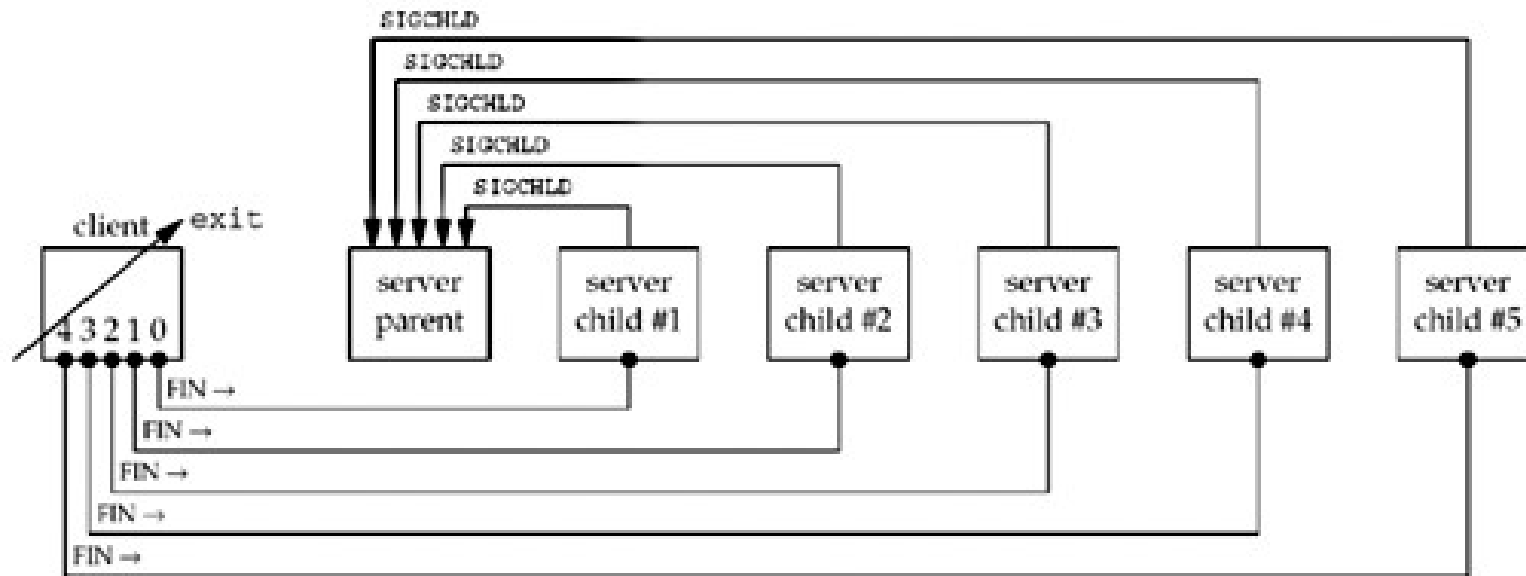


`wait()`



- Create 5 connections from a client to a forking server
- When the client terminates, all open descriptors are closed automatically by the kernel → five connections ended simultaneous

# waitpid()



- Client terminates, closing all five connections, terminating all five children → four children are zombies
- It can happen when many users connect to a server
- → we have to use *waitpid()*

## waitpid()

pid\_t waitpid (pid\_t pid, int \*statloc, int options);

- $pid < 0$ : wait for any child process whose process group ID is equal to the absolute value of *pid*.
- **pid = -1**: wait for any child process.
- $pid = 0$ : wait for any child process whose process group ID is equal to that of the calling process
- $pid > 0$ : wait for the child whose process ID is equal to the value of *pid*
- Without option WNOHANG, waitpid blocks until the status change
- With option WNOHANG, waitpid returns immediately
- Return
  - Pid of the child whose state has changed
  - with option WNOHANG, return 0 if the specified process has not changed status.

## void sig\_chld(int signo)

```
void sig_chld(int signo)
{
    pid_t pid;
    int stat;
    pid = waitpid(-1, &stat, WNOHANG );
    printf("child %d terminated\\", pid);
}
```

- **WNOHANG:** `waitpid()` does not block
- **while loop:** `waitpid()` repeatedly until there is no child process change status, i.e until `waitpid` returns 0.

# Forking server

```
pid_t pid;
int listenfd, connfd;
//Step 1: Construct socket
//Step 2: Bind address to socket
//Step 3: Listen request from client

// wait for a child process to stop
signal(SIGCHLD, sig_chld);
//Step 4: Communicate with client
while (1) {
    connfd = accept (listenfd, ... );
    if( (pid = fork()) == 0) {// process in child
        close(listenfd); // child closes listening socket
        doit(connfd);    // process the request
        close(connfd);   // done with this client
        exit(0);         // child terminates
    }
    close(connfd); // parent closes connected socket
}
```

# Handling EINTR Errors

- When a process is blocked in a **slow system call** and the process catches a signal and the signal handler returns, the system call can return an error of EINTR.
- Slow system call: connect, accept, send, recv...
- Not all kernels automatically restart some interrupted system calls
- We must rewrite function to handle EINTR error

```
while(1)
    if ((connfd = accept (listenfd...) < 0) {
        if (errno == EINTR)
            continue;           /* back to for () */
        else
            perror ("Error: ");
    }
```

## Other problems

- Connection abort before *accept* return
- Termination of server process
- Crashing of sever host
- Crashing and Reboot of server host



# MULTI-THREAD SERVER

---



# pthread\_create()

```
#include <pthread.h>
int pthread_create(pthread_t *tid, const pthread_attr_t *attr,
                  void *(*routine) (void *), void *arg);
```

- Create a new thread
- Parameters:
  - [OUT] `tid`: points to ID of the new thread
  - [IN] `attr`: points to structure whose contents are used to determine attributes for the new thread
  - [IN] `routine`: the new thread starts execution by invoking `routine()`
  - [IN] `arg`: points to the argument is passed as the sole argument of `routine()`
- Return:
  - On success, returns 0
  - On error, returns an error number
- **Compile and link with *-pthread***

# pthread\_create()

- By default, the new thread is joinable:
  - Not automatically cleaned up by GNU/Linux when it terminates
  - the thread's exit state hangs around in the system until another thread calls `pthread_join()` to obtain its return value
- Detached thread is cleaned up automatically when it terminates
  - Another thread may not obtain its return value
- Detach a thread: `int pthread_detach(pthread_t tid)`
  - On success, returns 0
  - On error, returns an error number

# Multi-thread TCP Echo Server

```
pthread_t tid;
int listenfd, connfd;
//Step 1: Construct socket
//Step 2: Bind address to socket
//Step 3: Listen request from client

//Step 4: Communicate with client
while (1) {
    connfd = accept (listenfd, ... );
    pthread_create(&tid, NULL, &client_handler,
                  (void *) connfd);
}
close(listenfd);
return 0;
```

# Multi-thread TCP Echo Server(cont.)

```
void *client_handler(void *arg) {
    int clientfd;
    int sendBytes, rcvBytes;
    char buff[BUFF_SIZE];

    pthread_detach(pthread_self());
    clientfd = (int) arg;
    while(1) {
        rcvBytes = recv(clientfd, buff, BUFF_SIZE, 0);
        if (rcvBytes < 0) {
            perror("\nError: ");
            break;
        }
        sendBytes = send(clientfd, buff, rcvBytes, 0);
        if (sendBytes < 0) {
            printf("\nError:");
            break;
        }
    }
    close(clientfd);
}
```

# Synchronize threads

- Since multiple threads can be running concurrently, accessing the shared variables:
  - The order of the accessing shared memory is unpredictable, so
  - The processing flow of the thread may be uncontrollable, and/or
  - The process crash
- Synchronize threads so that only one thread can access shared memory:
  - Inter-lock
  - Semaphore
  - Mutex

# Mutex

```
#include <pthread.h>
int pthread_mutex_lock(pthread_mutex_t * mptr);
int pthread_mutex_unlock(pthread_mutex_t * mptr);
```

- The thread can access the shared variable only when it hold the mutex
- `pthread_mutex_lock()`: lock a mutex
- `pthread_mutex_unlock()` : unlock a mutex
- If the thread try to lock a mutex that is already locked by some other thread, it is blocked until the mutex is unlocked.

```
void *routine(void *arg) {
    //...
    pthread_mutex_lock(mptr);
    // access shared memory
    pthread_mutex_unlock(mptr);
    //...
}
```

# fork() VS pthread\_create()

## fork()

- Heavy-weight
- Passing information from the parent to the child before the fork is easy
- Returning information from the child to the parent takes more work
- Needn't synchronize processes
- Greater isolation between the parent and the child

## pthread\_create()

- Light-weight
- Passing information from a thread to the others is easy
- Don't need signal-driven processing when the threads ends.
- May synchronize threads
- If a thread crashes, process may crash